

Random Testing

Question 1. Consider the following concurrent program with two threads and shared variable `x`. Variables `tmp1` and `tmp2` are local to the respective threads. This program has a concurrency bug: it can lose an update to `x`.

Thread 1	Thread 2
1: <code>tmp1=x</code>	4: <code>tmp2=x</code>
2: <code>tmp1=tmp1+1</code>	5: <code>tmp2=tmp2+1</code>
3: <code>x=tmp1</code>	6: <code>x=tmp2</code>

a. Write one possible execution of the six statements that does not cause a concurrency bug.

Answer: Any order which performs a read and a write before the second read, e.g., 1 2 3 4 5 6
OR 4 5 6 1 2 3.

b. Write one possible execution of the six statements that does trigger a concurrency bug.

Answer: Any order which performs both reads before any write, e.g., 1 4 2 5 3 6.

c. What is the depth of the concurrency bug?

Answer: 2

d. Specify the ordering constraints needed to trigger the bug.

Answer: (4, 3) (1, 6)

Question 2. Consider the following pseudo-Java function, in which `HashMap<char, int>` is used. A `HashMap<K, V>` is a data structure that associates a value of type `V` to a key of type `K`. The value `v` associated with a key `k` can be set with the API call `put(k, v)`, and the value associated with the key `k` is returned by the API call `get(k)`. For this problem, if no value has been associated with `k`, then assume `get(k)` returns 0.

```
double charRatio(String s, char a, char b) {
    int N = s.length();
    HashMap<char, int> counts = new HashMap<char, int>();
    for (int i = 0; i < N; i++) {
        char c = s.charAt(i);
        int v = counts.get(c);
        counts.put(c, v+1);
    }
    return counts.get(a) / counts.get(b);
}
```

```
}
```

Describe how you could use a fuzzer to test this function. What bugs would you expect a fuzzer to identify in this function? What bugs would be more challenging for a fuzzer to identify? Explain your reasoning fully, including any assumptions you are making.

Answer: There are three main bugs in this program:

- the possibility of a null dereference if `s == null`,
- the possibility of a division by zero when `b` doesn't appear in the input string, and
- the possibility of `counts.get(a) / counts.get(b)` not equaling the correct ratio of `a`'s to `b`'s in the input string due to integer division.

Fuzzing would likely quickly detect the division-by-zero error by generating a string `s` with no instances of the char `b`. This would depend on the implementation of the fuzzer. If the fuzzer only generated '0' and '1's, then it would likely be difficult for `b` not to appear in `s`. On the other hand, if the fuzzer uniformly generated legal strings of length 100 from the ASCII character set (and `b` were uniformly chosen from the ASCII character set), then there would be a 45% chance of the bug being triggered. (As the length of the input string grows, the probability of selecting `b` so that `b` were not in `s` would vanish quickly, though.)

The null-dereference error would also likely be caught by the fuzzer using a similar argument.

The integer division bug would be harder to detect, as fuzzing does not generally entail matching the output of a function against an expected output: we usually just give the function random strings until we detect a crash.